Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn

There will be some functions that start with the word "grader" ex: grader_weights(), grader_sigmoid(), grader_logloss() etc, you should not change those function definition.

Every Grader function has to return True.

Importing packages

In [1]:
```
import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model
```

Creating custom dataset

In [2]:
```
# please don't change random_state

X, y = make_classification(n_samples = 50000, n_features = 15, n_informative = 10, n_redundant = 5,
                           n_classes = 2, weights = [0.7], class_sep = 0.7, random_state = 15)

# make_classification is used to create custom dataset
# https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make_classification.html
```

In [3]:
```
X.shape, y.shape
```

Out[3]:
```
((50000, 15), (50000,))
```

Splitting data into train and test

In [4]:
```
#please don't change random state
# you need not standardize the data as it is already standardized

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 15)
```

In [5]:
```
X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

Out[5]:
```
((37500, 15), (37500,), (12500, 15), (12500,))
```

SGD classifier

In [6]:
```
# alpha : float
# Constant that multiplies the regularization term.

# eta0 : double
# The initial learning rate for the ‘constant’, ‘invscaling’ or ‘adaptive’ schedules.

clf = linear_model.SGDClassifier(eta0 = 0.0001, alpha = 0.0001, loss = 'log', random_state = 15,
                                 penalty = 'l2', tol = 1e-3, verbose = 2, learning_rate = 'constant')
clf
#https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
```

Out[6]:
```
SGDClassifier(eta0=0.0001, learning_rate='constant', loss='log',
              random_state=15, verbose=2)
```

In [7]:
```
clf.fit(X = X_train, y = y_train) # fitting our model
```

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.01 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.01 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.02 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.03 seconds.
-- Epoch 5
```

```
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.04 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.04 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.05 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.06 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.07 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.07 seconds.
Convergence after 10 epochs took 0.07 seconds
```

Out[7]:
```
SGDClassifier(eta0=0.0001, learning_rate='constant', loss='log',
              random_state=15, verbose=2)
```

In [8]:
```python
#clf.coef_ will return the weights
#clf.coef_.shape will return the shape of weights
#clf.intercept_ will return the intercept term

clf.coef_, clf.coef_.shape, clf.intercept_
```

Out[8]:
```
(array([[-0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
          0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.18084126,
          0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
 (1, 15),
 array([-0.8531383]))
```

Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

1. We will be giving you some functions, please write code in that functions only.
2. After every function, we will be giving you expected output, please make sure that you get that output.

- Initialize the weight_vector and intercept term to zeros (Write your code in def initialize_weights())
- Create a loss function (Write your code in def logloss())

  $log loss = -1*\frac{1}{n}\Sigma_{for each Yt,Y_{pred}}(Ytlog10(Y_{pred})+(1-Yt)log10(1-Y_{pred}))$

- for each epoch:
  - for each batch of data points in train: (keep batch size=1)
    - calculate the gradient of loss function w.r.t each weight in weight vector (write your code in def gradient_dw())

      $dw^{(t)} = x_n(y_n − σ((w^{(t)})^{T} x_n+b^{t}))- \frac{λ}{N}w^{(t)})$

    - Calculate the gradient of the intercept (write your code in def gradient_db()) check this

      $ db^{(t)} = y_n- σ((w^{(t)})^{T} x_n+b^{t}))$

    - Update weights and intercept (check the equation number 32 in the above mentioned pdf):
      $w^{(t+1)}← w^{(t)}+α(dw^{(t)}) $

      $b^{(t+1)}←b^{(t)}+α(db^{(t)}) $
  - calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
  - And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training
  - append this loss in the list (this will be used to see how loss is changing for each epoch after the training is over)

Initialize weights

In [9]:
```python
#you use zeros_like function to initialize zero,
#https://numpy.org/doc/stable/reference/generated/numpy.zeros_like.html
#initialize bias to zero

def initialize_weights(row_vector):
```

```
    ''' In this function, we will initialize our weights and bias'''

    w = np.zeros_like(row_vector)
    b = 0

    return w,b
```

In [10]:

```
dim = X_train[0]

# w,b = initialize_weights(row_vector)
w,b = initialize_weights(dim)

print('w =',(w))
print('b =',str(b))
```

```
w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
b = 0
```

Grader function - 1

In [11]:

```
dim = X_train[0]
w,b = initialize_weights(dim)

def grader_weights(w,b):
    assert((len(w) == len(dim)) and b == 0 and np.sum(w) == 0.0)
    return True

grader_weights(w,b)
```

Out[11]: True

Compute sigmoid

$sigmoid(z)= 1/(1+exp(-z))$

In [12]:

```
# https://numpy.org/doc/stable/reference/generated/numpy.exp.html

def sigmoid(z):
    ''' In this function, we will return sigmoid of z'''

    # compute sigmoid(z) and return
    sigmoid_ = (1 / (1 + np.exp(-z)))

    return sigmoid_
```

Grader function - 2

In [13]:

```
def grader_sigmoid(z):

    val = sigmoid(z)
    assert(val == 0.8807970779778823)

    return True

grader_sigmoid(2)
```

Out[13]: True

Compute loss

$log loss = -1*\frac{1}{n}\Sigma_{for each Yt,Y_{pred}}(Ytlog10(Y_{pred})+(1-Yt)log10(1-Y_{pred}))$

In [14]:

```
# https://numpy.org/doc/stable/reference/generated/numpy.log10.html

def logloss(y_true, y_pred):

    sum_ = 0

    for i in range(len(y_true)):
        sum_ += y_true[i] * np.log10(y_pred[i]) + (1- y_true[i]) * (np.log10(1-y_pred[i]))

    loss = -1 * (1/len(y_true)) * sum_

    return loss
```

Grader function - 3

In [15]:

```
#round off the value to 8 values

def grader_logloss(true, pred):

    loss = logloss(true,pred)
```

```
    assert(np.round(loss,6) == 0.076449)

    return True

true = np.array([1,1,0,1,0])
pred = np.array([0.9,0.8,0.1,0.8,0.2])

grader_logloss(true, pred)
```

Out[15]: True

Compute gradient w.r.to 'w'

$dw^{(t)} = x_n(y_n − σ((w^{(t)})^{T} x_n+b^{t}))- \frac{λ}{N}w^{(t)}$

In [16]:

```
# https://numpy.org/doc/stable/reference/generated/numpy.transpose.html

def gradient_dw(x, y, w, b, alpha, N):

    '''In this function, we will compute the gardient w.r.to w '''

    dw = x * (y - sigmoid(np.dot(np.transpose(w),x) + b)) - (( alpha / N) * w)

    return dw
```

Grader function - 4

In [17]:

```
def grader_dw(x, y, w, b, alpha, N):

    grad_dw = gradient_dw(x,y,w,b,alpha,N)
    assert(np.round(np.sum(grad_dw),5) == 4.75684)
    return True

grad_x = np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783286,
       -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
        3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])
grad_y = 0
grad_w = np.array([ 0.03364887,  0.03612727,  0.02786927,  0.08547455, -0.12870234,
       -0.02555288,  0.11858013,  0.13305576,  0.07310204,  0.15149245,
       -0.05708987, -0.064768  ,  0.18012332, -0.16880843, -0.27079877])

grad_b = 0.5
alpha = 0.0001
N = len(X_train)

grader_dw(grad_x, grad_y, grad_w, grad_b, alpha,N)
```

Out[17]: True

Compute gradient w.r.to 'b'

$ db^{(t)} = y_n- σ((w^{(t)})^{T} x_n+b^{t})$

In [18]:

```
def gradient_db(x, y, w, b):
    '''In this function, we will compute gradient w.r.to b '''

    db = y - (sigmoid(np.dot(np.transpose(w), x) + b))

    return db
```

Grader function - 5

In [19]:

```
def grader_db(x, y, w, b):

    grad_db=gradient_db(x,y,w,b)
    assert(np.round(grad_db,4) == -0.3714)
    return True

grad_x = np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783286,
       -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
        3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])

grad_y = 0.5
grad_b = 0.1

grad_w = np.array([ 0.03364887,  0.03612727,  0.02786927,  0.08547455, -0.12870234,
       -0.02555288,  0.11858013,  0.13305576,  0.07310204,  0.15149245,
       -0.05708987, -0.064768  ,  0.18012332, -0.16880843, -0.27079877])

alpha = 0.0001
N = len(X_train)

grader_db(grad_x, grad_y, grad_w, grad_b)
```

Out[19]: True

In [20]:
```
# prediction function used to compute predicted_y given the dataset X

def pred(w, b, X):

    N = len(X)
    predict = []

    for i in range(N):
        z = np.dot(w, X[i]) + b
        predict.append(sigmoid(z))

    return np.array(predict)
```

https://youtu.be/Zc6RBeTrYjE?t=503 : SGD Assignment def pred(w,b, x): N = len(x) predict = [] for i in range(N): z = np.dot(w,x[i]) + b if sigmoid(z) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+b))) predict.append(1) else: predict.append(0) return np.array(predict)

Implementing logistic regression

In [21]:
```
def train(X_train, y_train, X_test, y_test, epochs, alpha, eta0, N):
    ''' In this function, we will implement logistic regression'''

    train_loss = []
    test_loss = []

    w, b = initialize_weights(X_train[0])

    for epoch in range(epochs):

        for i in range(N):
            dw = gradient_dw(X_train[i], y_train[i], w, b, alpha, N)
            db = gradient_db(X_train[i], y_train[i], w, b)

#              w = w - (eta0 * dw)
#              b = b - (eta0 * db)

# https://rasbt.github.io/mlxtend/user_guide/general_concepts/gradient-optimization/
            w = w + (eta0 * dw)
            b = b + (eta0 * db)

        train_pred = pred(w, b, X_train)

        train_pred_loss = logloss(y_train, train_pred)
        train_loss.append(train_pred_loss)

        test_pred = pred(w, b, X_test)

        test_pred_loss = logloss(y_test,test_pred)
        test_loss.append(test_pred_loss)

    return w, b, train_loss, test_loss
```

In [22]:
```
# https://youtu.be/Zc6RBeTrYjE : SGD Assignment

# alpha = 0.001
# eta0 = 0.001
eta0 = 0.0001
alpha = 0.0001
N = len(X_train)
epochs = 20

w, b, train_loss, test_loss = train(X_train, y_train, X_test, y_test, epochs, alpha, eta0, N)
```

In [23]:
```
#print thr value of weights w and bias b

print(w)
print(b)
```

```
[-4.29394713e-01  1.92911531e-01 -1.48319226e-01  3.38095811e-01
 -2.20731189e-01  5.69669865e-01 -4.45186056e-01 -9.00099226e-02
  2.21598219e-01  1.73588003e-01  1.98538391e-01 -4.13172177e-04
 -8.11250040e-02  3.39070544e-01  2.29369069e-02]
-0.8897519322788823
```

In [24]:
```
# these are the results we got after we implemented sgd and found the optimal weights and intercept

w - clf.coef_, b - clf.intercept_
```

Out[24]:
```
(array([[-0.0060278 ,  0.00743588,  0.00027113, -0.00334826, -0.01254449,
          0.00950408,  0.00723877,  0.00407821,  0.01232502, -0.00725326,
          0.00148649, -0.00463233, -0.00152131,  0.00054253,  0.0002697 ]]),
 array([-0.03661363]))
```

Goal of assignment

Compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in order of 10^-2

Grader function - 6

In [25]:

```
#this grader function should return True
#the difference between custom weights and clf.coef_ should be less than or equal to 0.05

def differece_check_grader(w, b, coef, intercept):

    val_array = np.abs(np.array(w - coef))
    assert(np.all(val_array <= 0.05))

    print('The custom weights are correct')
    return True

differece_check_grader(w, b, clf.coef_, clf.intercept_)
```

```
The custom weights are correct
```

Out[25]:

```
True
```

Plot your train and test loss vs epochs

plot epoch number on X-axis and loss on Y-axis and make sure that the curve is converging

In [26]:

```
from matplotlib import pyplot as plt
epoch = np.arange(20) + 1

plt.figure(figsize = (12,5))
plt.plot(epoch,train_loss , label='Train Log Loss')
plt.plot(epoch,test_loss, label='Test Log Loss')
plt.xticks(epoch)
plt.xlabel("Epoch number")
plt.ylabel("Log Loss")
plt.legend()
plt.show()
```