

Transfer Learning Assignment

Download all the data in this [rar_file](#) , it contains all the data required for the assignment. When you unrar the file you'll get the files in the following format: **path/to/the/image.tif,category**

where the categories are numbered 0 to 15, in the following order:

```
0 letter
1 form
2 email
3 handwritten
4 advertisement
5 scientific report
6 scientific publication
7 specification
8 file folder
9 news article
10 budget
11 invoice
12 presentation
13 questionnaire
14 resume
15 memo
```

There is a file named as 'labels_final.csv' , it consists of two columns. First column is path which is the required path to the images and second is the class label.

The dataset that you are dealing with is quite large 3.7 GB and hence there are two methods to import the data to Colab

Method 1 - You can use gdown module to get the data directly from Google drive to Colab

The syntax is as follows,

```
! gdown --id file_id
```

for ex - running the below cell will

```
import the rvl-cdip.rar dataset
! gdown --id 1Z4TyI7FcFVEx8qd14j09qxvxaqLSqoEu
```

Method 2 - You can also import the data using wget function

<https://www.youtube.com/watch?v=BPUfVq7RaY8>

Unzip `rar` file using,

```
#unrar the file
get_ipython().system_raw("unrar x rvl-cdip.rar")
```

2. On this image data, you have to train 3 types of models as given below You have to split the data into Train and Validation data.

```
In [1]: import os
import numpy as np
import pandas as pd

from datetime import datetime

import seaborn as sns
import matplotlib.pyplot as plt
plt.style.use('fivethirtyeight')

import tensorflow
import tensorflow as tf
from tensorflow.keras import Model
from tensorflow.keras import Input

from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Conv2D
from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Flatten
from tensorflow.keras.layers import MaxPool2D
from tensorflow.keras.layers import Activation
from tensorflow.keras.callbacks import Callback
```

```

from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.callbacks import TensorBoard
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.callbacks import ReduceLR0nPlateau

from tensorflow.keras.utils import plot_model
from tensorflow.keras.applications import VGG16
from tensorflow.keras.preprocessing.image import ImageDataGenerator
%load_ext tensorboard

```

```

In [2]: df = pd.read_csv('labels_final.csv', dtype = str)

```

```

In [3]: print(df.shape)
df.head()

```

```

(48000, 2)

```

```

Out[3]:

```

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	3
1	imagesl/l/x/t/lxt19d00/502213303.tif	3
2	imagesx/x/e/d/xed05a00/2075325674.tif	2
3	imageso/o/j/b/ojb60d00/517511301+-1301.tif	3
4	imagesq/q/z/k/qzk17e00/2031320195.tif	7

1. Try not to load all the images into memory, use the generators that we have given the reference notebooks to load the batch of images only during the train data. or you can use this method also <https://medium.com/@vijayabhaskar96/tutorial-on-keras-imagedatagenerator-with-flow-from-dataframe-8bd5776e45c1>

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c>

Note- In the reference notebook you were dealing with jpg images, in the given dataset you are dealing with tiff images. Imagedatagenrator works with both type of images. If you want to use custom data pipeline then you have to convert your tiff images to jpg images.

1. You are free to choose Learning rate, optimizer, loss function, image augmentation, any hyperparameters. but you have to use the same architecture what we are asking below.
2. Use tensorboard for every model and analyse your gradients. (you need to upload the screenshots for each model for evaluation)
1. You can check about Transfer Learning in this link - <https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>

<https://www.appliedaicourse.com/lecture/11/applied-machine-learning-online-course/3426/code-example-cats-vs-dogs/8/module-8-neural-networks-computer-vision-and-deep-learning>

1. Do print model.summary() and draw model_plots for each of the model.

Threshold Accuracies

Model	Accuracy
Model 1	60 %
Model 2	60 %
Model 3	6 %

Model-1

1. Use VGG-16 pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, add a new Conv block (1 Conv layer and 1 Maxpooling), 2 FC layers and an output layer to classify 16 classes. You are free to choose any

- hyperparameters/parameters of conv block, FC layers, output layer.
3. Final architecture will be **INPUT --> VGG-16 without Top Layers(FC) --> Conv Layer --> Maxpool Layer --> 2 FC layers --> Output Layer**
4. Print model.summary() and plot the architecture of the model.
- [Reference for plotting model](#)
5. Train only new Conv block, FC layers, output layer. Don't train the VGG-16 network.

```
In [4]: # https://stackoverflow.com/a/35765856

# !watch -n1 nvidia-smi
!nvidia-smi
```

Thu Jul 14 14:58:15 2022

```
+-----+
| NVIDIA-SMI 440.33.01      Driver Version: 440.33.01      CUDA Version: 11.0      |
+-----+-----+-----+-----+-----+-----+
| GPU  Name            Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp   Perf      Pwr:Usage/Cap|  Memory-Usage | GPU-Util  Compute M. |
+-----+-----+-----+-----+-----+-----+
|    0   Tesla T4              On   | 00000000:00:1E:0 Off |            0         |
| N/A   33C    P8         8W / 70W |  0MiB / 15109MiB |      0%    Default   |
+-----+-----+-----+-----+-----+-----+

+-----+-----+-----+-----+-----+-----+
| Processes:                                                       GPU Memory |
|  GPU       PID    Type    Process name                       Usage    |
+-----+-----+-----+-----+-----+-----+
| No running processes found                                     |
+-----+-----+-----+-----+-----+-----+
```

```
In [5]: os.mkdir('Model_Weights')
os.mkdir('Model_Plots')
```

```
In [6]: tensorflow.random.set_seed(42)
```

```
In [7]: # https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator

datagen = ImageDataGenerator(rescale = 1./255, shear_range = 0.2, zoom_range = 0.2,
                             validation_split = 0.25)
```

```
In [8]: # https://medium.datadriveninvestor.com/keras-imagedatagenerator-methods-an-easy-guide-550ecd3c0a92

BATCH_SIZE = 32
DATA_DIR = 'data_final'

train_gen = datagen.flow_from_dataframe(dataframe = df, directory = DATA_DIR,
                                       x_col = 'path', y_col = 'label', target_size = (256,256),
                                       class_mode = 'categorical', batch_size = BATCH_SIZE,
                                       subset = 'training', seed = 7, shuffle = True, color_mode = 'rgb')

test_gen = datagen.flow_from_dataframe(dataframe = df, directory = DATA_DIR,
                                       x_col = 'path', y_col = 'label', target_size = (256,256),
                                       class_mode = 'categorical', batch_size = BATCH_SIZE,
                                       subset = 'validation', seed = 7, shuffle = True, color_mode = 'rgb')
```

Found 36000 validated image filenames belonging to 16 classes.
Found 12000 validated image filenames belonging to 16 classes.

```
In [9]: base_model = VGG16(weights = 'imagenet', input_shape = (256, 256, 3), include_top = False)
base_model.summary()
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58892288/58889256 [=====] - 0s 0us/step
Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928

block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

INPUT --> VGG-16 without Top layers(FC) --> Conv Layer --> Maxpool Layer --> 2 FC layers --> Output Layer

In [10]:

```
tensorflow.keras.backend.clear_session()

# https://keras.io/guides/transfer_learning/
# https://catalog.ngc.nvidia.com/orgs/nvidia/containers/dli-dl-fundamentals

base_model.trainable = False
inputs = Input(shape = (256, 256, 3))

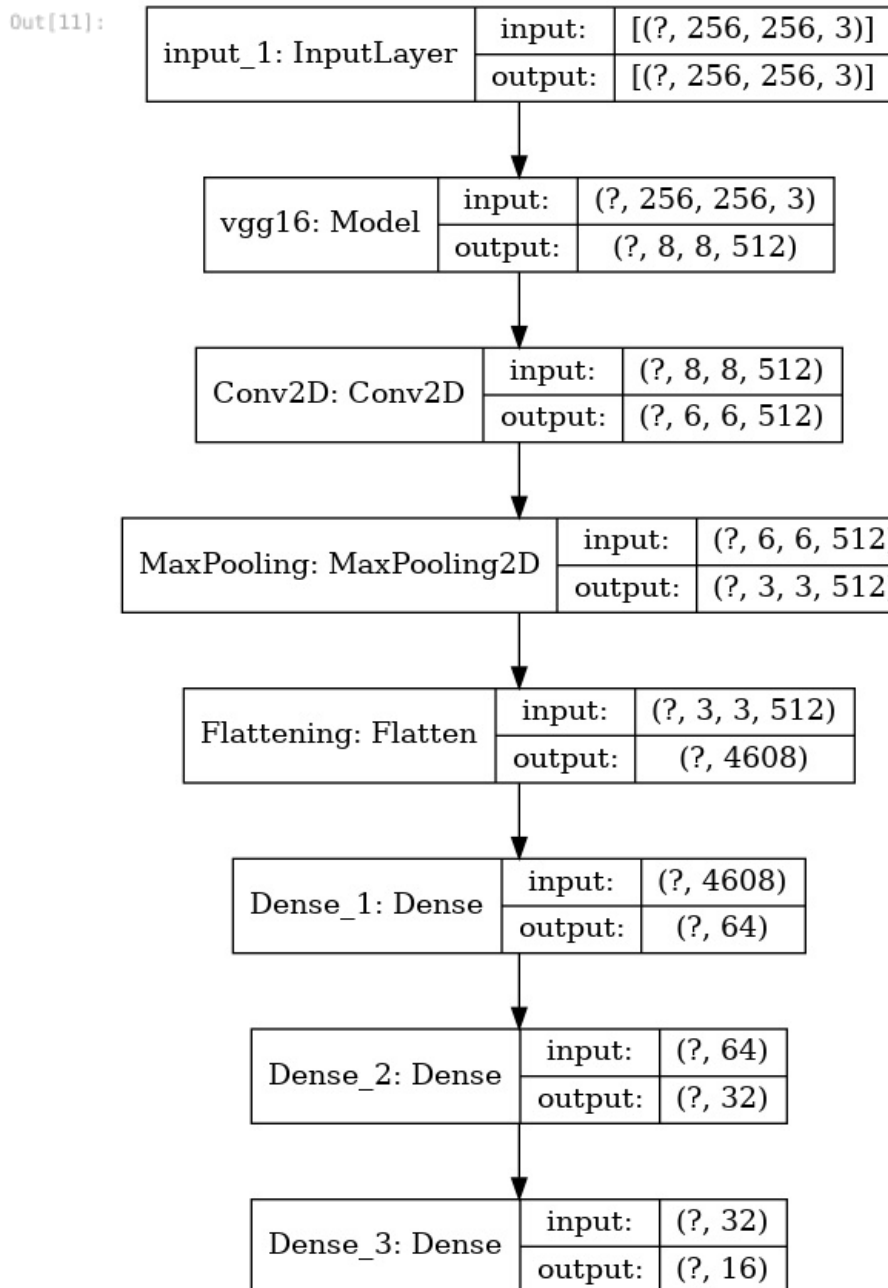
x = base_model(inputs, training = False)
x = Conv2D(filters = 512, kernel_size = (3, 3), activation = 'relu', name = 'Conv2D')(x)
x = MaxPool2D(name = 'MaxPooling')(x)
x = Flatten(name = 'Flattening')(x)
x = Dense(64, activation = 'relu', name = 'Dense_1')(x)
x = Dense(32, activation = 'relu', name = 'Dense_2')(x)
output_ = Dense(16, activation = 'softmax', name = 'Dense_3')(x)

final_model_1 = Model(inputs = inputs, outputs = output_, name = 'Model_1')
final_model_1.summary()
```

Model: "Model_1"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
vgg16 (Model)	(None, 8, 8, 512)	14714688
Conv2D (Conv2D)	(None, 6, 6, 512)	2359808
MaxPooling (MaxPooling2D)	(None, 3, 3, 512)	0
Flattening (Flatten)	(None, 4608)	0
Dense_1 (Dense)	(None, 64)	294976
Dense_2 (Dense)	(None, 32)	2080
Dense_3 (Dense)	(None, 16)	528
=====		
Total params: 17,372,080		
Trainable params: 2,657,392		
Non-trainable params: 14,714,688		

```
In [11]: plot_model(final_model_1, to_file = 'Model_Plots/Model_1.png', show_shapes = True)
```



```
In [12]: final_model_1.compile(optimizer = 'adam', loss = 'categorical_crossentropy',
                               metrics = ['accuracy'])
```

```
In [13]: logdir = 'logs/model_1/' + datetime.now().strftime('%Y%m%d_%H%M%S')
tensorBoard = TensorBoard(log_dir = logdir, histogram_freq = 1)

filepath = 'modelCheck/model_1/epo_{epoch:02d}-accu_{val_accuracy:.4f}.hdf5'
modelCheck = ModelCheckpoint(filepath, monitor = 'val_accuracy', verbose = 1)

earlyStopping = EarlyStopping(monitor = 'accuracy', patience = 2, verbose = 1)
reduceLr = ReduceLROnPlateau(monitor = 'accuracy', factor = 0.9, patience = 2, verbose = 1)

callBack_List = [tensorBoard, modelCheck, earlyStopping, reduceLr]
```

```
In [14]: step_train = train_gen.n // train_gen.batch_size
step_test = test_gen.n // test_gen.batch_size
```

```
In [15]: EPOCHS_1 = 6
model_1 = final_model_1.fit(train_gen, steps_per_epoch = step_train, epochs = EPOCHS_1,
                             validation_data = test_gen, validation_steps = step_test, callbacks = callBack_List)
```

```
Epoch 1/6
1125/1125 [=====] - ETA: 0s - loss: 1.6810 - accuracy: 0.4734
Epoch 00001: saving model to modelCheck/model_1/epo_01-accu_0.5727.hdf5
1125/1125 [=====] - 713s 634ms/step - loss: 1.6810 - accuracy: 0.4734 - val_loss: 1.3635
```

```

- val_accuracy: 0.5727 - lr: 0.0010
Epoch 2/6
1125/1125 [=====] - ETA: 0s - loss: 1.2388 - accuracy: 0.6151
Epoch 00002: saving model to modelCheck/model_1/epo_02-accu_0.6342.hdf5
1125/1125 [=====] - 671s 596ms/step - loss: 1.2388 - accuracy: 0.6151 - val_loss: 1.2049
- val_accuracy: 0.6342 - lr: 0.0010
Epoch 3/6
1125/1125 [=====] - ETA: 0s - loss: 1.0989 - accuracy: 0.6604
Epoch 00003: saving model to modelCheck/model_1/epo_03-accu_0.6637.hdf5
1125/1125 [=====] - 673s 598ms/step - loss: 1.0989 - accuracy: 0.6604 - val_loss: 1.1289
- val_accuracy: 0.6637 - lr: 0.0010
Epoch 4/6
1125/1125 [=====] - ETA: 0s - loss: 1.0028 - accuracy: 0.6936
Epoch 00004: saving model to modelCheck/model_1/epo_04-accu_0.6810.hdf5
1125/1125 [=====] - 673s 598ms/step - loss: 1.0028 - accuracy: 0.6936 - val_loss: 1.0616
- val_accuracy: 0.6810 - lr: 0.0010
Epoch 5/6
1125/1125 [=====] - ETA: 0s - loss: 0.9284 - accuracy: 0.7145
Epoch 00005: saving model to modelCheck/model_1/epo_05-accu_0.6883.hdf5
1125/1125 [=====] - 676s 601ms/step - loss: 0.9284 - accuracy: 0.7145 - val_loss: 1.0538
- val_accuracy: 0.6883 - lr: 0.0010
Epoch 6/6
1125/1125 [=====] - ETA: 0s - loss: 0.8707 - accuracy: 0.7345
Epoch 00006: saving model to modelCheck/model_1/epo_06-accu_0.7028.hdf5
1125/1125 [=====] - 675s 600ms/step - loss: 0.8707 - accuracy: 0.7345 - val_loss: 0.9977
- val_accuracy: 0.7028 - lr: 0.0010

```

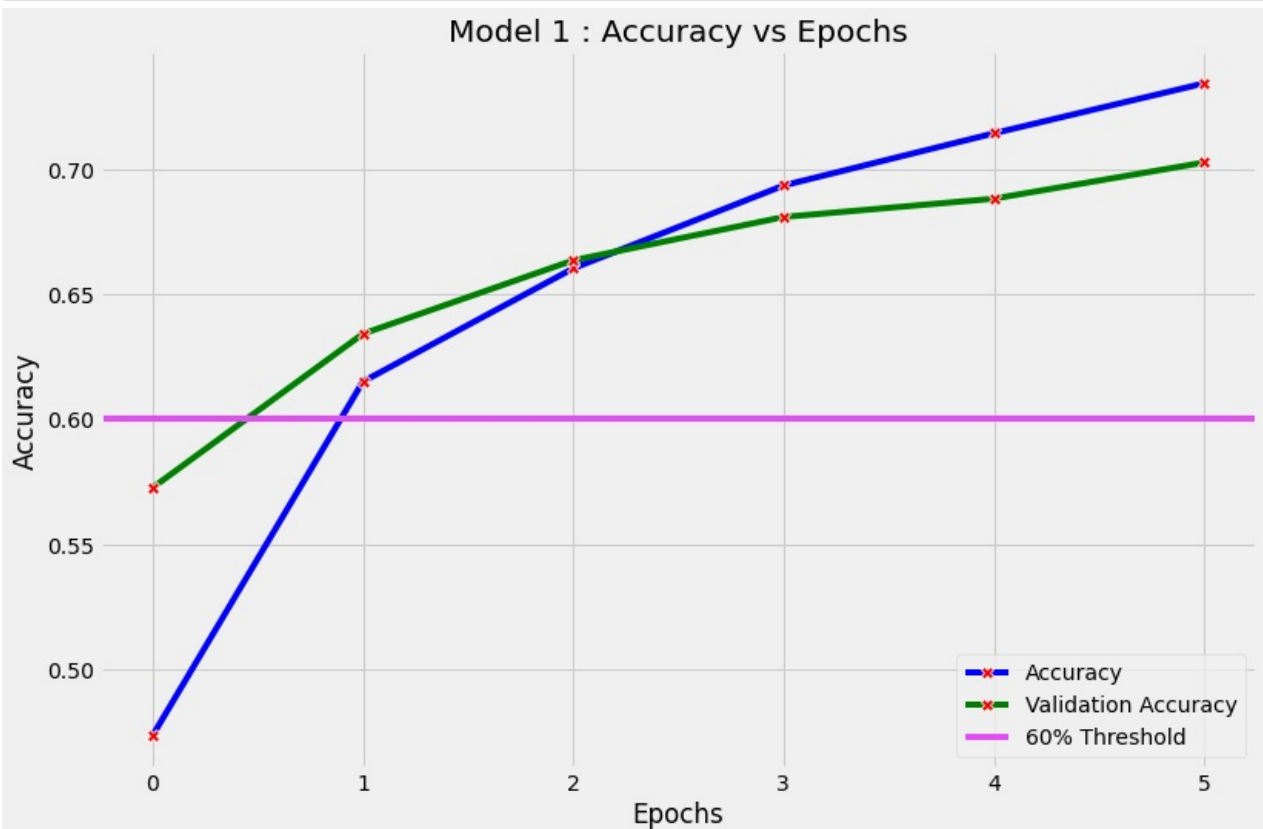
```
In [16]: final_model_1.save_weights('Model_Weights/best_model_1.h5')
```

```
In [17]: # https://wellsr.com/python/seaborn-line-plot-data-visualization/
# https://seaborn.pydata.org/generated/seaborn.lineplot.html

epochs_ = list(range(len(model_1.history['accuracy'])))
plt.figure(figsize = (12, 8))

sns.lineplot(y = 'accuracy', data = model_1.history, x = epochs_, label = 'Accuracy',
             color = 'b', marker = 'X', mfc = 'red', ms = 8)
sns.lineplot(y = 'val_accuracy', data = model_1.history, x = epochs_,
             label = 'Validation Accuracy', color = 'g', marker = 'X', mfc = 'red', ms = 8)

plt.axhline(0.6, color = '#E44CF6', label = '60% Threshold')
plt.title('Model 1 : Accuracy vs Epochs') ; plt.legend(loc = 4)
plt.xticks(epochs_) ; plt.xlabel('Epochs') ; plt.ylabel('Accuracy') ; plt.show()
```



Observation

- Only around 15% of parameters are trainable in total parameters.

- With 6 epochs we got ~73.45% accuracy and the loss got reduced from 1.6810 to 0.8707.
- The loss and accuracy shows that with higher epoch we can reduce loss and increase the accuracy of the model.

Model-2

1. Use [VGG-16](#) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, don't use FC layers, use conv layers only as Fully connected layer. Any FC layer can be converted to a CONV layer. This conversion will reduce the No of Trainable parameters in FC layers.
For example, an FC layer with K=4096 that is looking at some input volume of size 7×7×512 can be equivalently expressed as a CONV layer with F=7,P=0,S=1,K=4096.
In other words, we are setting the filter size to be exactly the size of the input volume, and hence the output will simply be 1×1×4096 since only a single depth column “fits” across the input volume, giving identical result as the initial FC layer. You can refer [this](#) link to better understanding of using Conv layer in place of fully connected layers.
3. Final architecture will be VGG-16 without FC layers(without top), 2 Conv layers identical to FC layers, 1 output layer for 16 class classification. **INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC -->Output Layer**
4. Print model.summary() and plot the architecture of the model.
[Reference for plotting model](#)
5. Train only last 2 Conv layers identical to FC layers, 1 output layer. Don't train the VGG-16 network.

INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC -->Output Layer

- $$W_{out} = \frac{W - F + 2P}{S} + 1$$
- $$H_{out} = \frac{H - F + 2P}{S} + 1$$

```
In [18]: tensorflow.keras.backend.clear_session()
```

```
In [19]: # https://keras.io/guides/transfer_learning/
# https://catalog.ngc.nvidia.com/orgs/nvidia/containers/dli-dl-fundamentals
# https://mein2work.medium.com/convertng-fc-layers-to-conv-layers-8a43880a44ed

base_model.trainable = False
inputs = Input(shape = (256, 256, 3))

x = base_model(inputs, training = False)
x = Conv2D(filters = 512, kernel_size = (8, 8), activation = 'relu')(x)
x = Conv2D(filters = 128, kernel_size = (1, 1), activation = 'relu')(x)
# x = Conv2D(filters = 16, kernel_size = (1, 1), activation = 'relu')(x)
x = tf.keras.layers.GlobalAveragePooling2D()(x)
output_ = Dense(16, activation='softmax')(x)

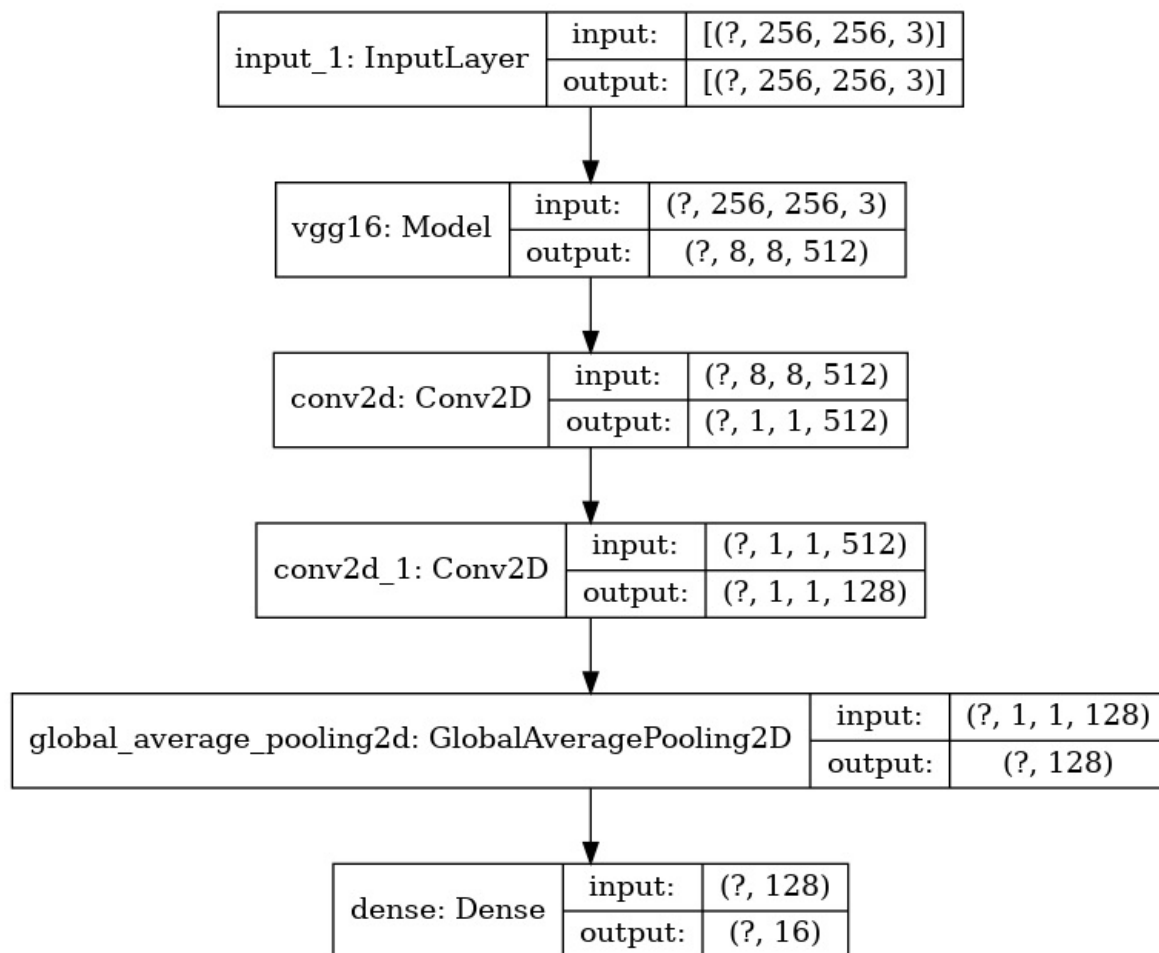
final_model_2 = Model(inputs = inputs, outputs = output_, name = 'Model_2')
final_model_2.summary()
```

Model: "Model_2"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
vgg16 (Model)	(None, 8, 8, 512)	14714688
conv2d (Conv2D)	(None, 1, 1, 512)	16777728
conv2d_1 (Conv2D)	(None, 1, 1, 128)	65664
global_average_pooling2d (Gl	(None, 128)	0
dense (Dense)	(None, 16)	2064
=====		
Total params: 31,560,144		
Trainable params: 16,845,456		
Non-trainable params: 14,714,688		

```
In [20]: plot_model(final_model_2, to_file = 'Model_Plots/Model_2.png', show_shapes = True)
```

```
Out[20]:
```



```
In [21]: logdir = 'logs/model_2/' + datetime.now().strftime('%Y%m%d_%H%M%S')
tensorBoard = TensorBoard(log_dir = logdir, histogram_freq = 1)

filepath = 'modelCheck/model_2/epo_{epoch:02d}-accu_{val_accuracy:.4f}.hdf5'
modelCheck = ModelCheckpoint(filepath, monitor = 'val_accuracy', verbose = 1)

callBack_List = [tensorBoard, modelCheck, earlyStopping, reduceLr]
```

```
In [22]: final_model_2.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])
```

```
In [23]: EPOCHS_2 = 6
model_2 = final_model_2.fit(train_gen, steps_per_epoch = step_train, epochs = EPOCHS_2,
                             validation_data = test_gen, validation_steps = step_test, callbacks = callBack_List)
```

```
Epoch 1/6
1125/1125 [=====] - ETA: 0s - loss: 1.4232 - accuracy: 0.5759
Epoch 00001: saving model to modelCheck/model_2/epo_01-accu_0.6572.hdf5
1125/1125 [=====] - 677s 602ms/step - loss: 1.4232 - accuracy: 0.5759 - val_loss: 1.1485
- val_accuracy: 0.6572 - lr: 0.0010
Epoch 2/6
1125/1125 [=====] - ETA: 0s - loss: 1.0631 - accuracy: 0.6741
Epoch 00002: saving model to modelCheck/model_2/epo_02-accu_0.6690.hdf5
1125/1125 [=====] - 675s 600ms/step - loss: 1.0631 - accuracy: 0.6741 - val_loss: 1.1056
- val_accuracy: 0.6690 - lr: 0.0010
Epoch 3/6
1125/1125 [=====] - ETA: 0s - loss: 0.9474 - accuracy: 0.7098
Epoch 00003: saving model to modelCheck/model_2/epo_03-accu_0.6873.hdf5
1125/1125 [=====] - 672s 598ms/step - loss: 0.9474 - accuracy: 0.7098 - val_loss: 1.0582
- val_accuracy: 0.6873 - lr: 0.0010
Epoch 4/6
1125/1125 [=====] - ETA: 0s - loss: 0.8784 - accuracy: 0.7284
Epoch 00004: saving model to modelCheck/model_2/epo_04-accu_0.7083.hdf5
1125/1125 [=====] - 673s 598ms/step - loss: 0.8784 - accuracy: 0.7284 - val_loss: 0.9910
- val_accuracy: 0.7083 - lr: 0.0010
Epoch 5/6
1125/1125 [=====] - ETA: 0s - loss: 0.8250 - accuracy: 0.7454
Epoch 00005: saving model to modelCheck/model_2/epo_05-accu_0.7057.hdf5
1125/1125 [=====] - 675s 600ms/step - loss: 0.8250 - accuracy: 0.7454 - val_loss: 0.9834
```



```
- val_accuracy: 0.7057 - lr: 0.0010
Epoch 6/6
1125/1125 [=====] - ETA: 0s - loss: 0.7710 - accuracy: 0.7603
Epoch 00006: saving model to modelCheck/model_2/epo_06-accu_0.7246.hdf5
1125/1125 [=====] - 674s 599ms/step - loss: 0.7710 - accuracy: 0.7603 - val_loss: 0.9490
- val_accuracy: 0.7246 - lr: 0.0010
```


```
In [24]: final_model_2.save_weights('Model_Weights/best_model_2.h5')
```

```
In [25]: # https://wellsr.com/python/seaborn-line-plot-data-visualization/
# https://seaborn.pydata.org/generated/seaborn.lineplot.html

epochs_ = list(range(len(model_2.history['accuracy'])))
plt.figure(figsize = (12, 8))

sns.lineplot(y = 'accuracy', data = model_2.history, x = epochs_, label = 'Accuracy',
             color = 'b', marker = 'X', mfc = 'red', ms = 8)
sns.lineplot(y = 'val_accuracy', data = model_2.history, x = epochs_,
             label = 'Validation Accuracy', color = 'g', marker = 'X', mfc = 'red', ms = 8)

plt.axhline(0.6, color = '#E44CF6', label = '60% Threshold')
plt.title('Model 2 : Accuracy vs Epochs') ; plt.legend(loc = 4)
plt.xticks(epochs_) ; plt.xlabel('Epochs') ; plt.ylabel('Accuracy') ; plt.show()
```



Observation

- Around 53% of parameters are trainable here.
- The loss was 0.7710 and accuracy was 76.03% in this model.
- This shows comparing to pervious model this model is giving much more better result in same 6 epochs.
- The curve also indicates that probably with more epoch we can increase the accuracy and reduce the loss.

Model-3

1. Use same network as Model-2 '**INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer**' and train only Last 6 Layers of VGG-16 network, 2 Conv layers identical to FC layers, 1 output layer.

```
In [26]: tensorflow.keras.backend.clear_session()
tensorflow.random.set_seed(42)
```

```
In [27]: # https://www.tensorflow.org/api_guides/python/random#Seeding
```

```
# https://www.w3resource.com/python/python-format.php
```

```
base_model = VGG16(weights = 'imagenet', input_shape = (256, 256, 3), include_top = False)

le_ = len(base_model.layers)

for layer in range(le_-6):
    base_model.layers[layer].trainable = False

count_true = 0
for layer in base_model.layers:
    if layer.trainable:
        count_true += 1
    print(f'{layer.name:15} {layer.trainable}')
print(f'\nNumber of Trainable Layers from base_model : {count_true}')
```

```
input_1      False
block1_conv1 False
block1_conv2 False
block1_pool  False
block2_conv1 False
block2_conv2 False
block2_pool  False
block3_conv1 False
block3_conv2 False
block3_conv3 False
block3_pool  False
block4_conv1 False
block4_conv2 False
block4_conv3 True
block4_pool  True
block5_conv1 True
block5_conv2 True
block5_conv3 True
block5_pool  True
```

Number of Trainable Layers from base_model : 6

In [28]: `base_model.summary()`

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0

=====
Total params: 14,714,688
Trainable params: 9,439,232
Non-trainable params: 5,275,456

```
In [29]: tensorflow.keras.backend.clear_session()

# https://keras.io/guides/transfer_learning/
# https://catalog.ngc.nvidia.com/orgs/nvidia/containers/dli-dl-fundamentals
# https://mein2work.medium.com/converting-fc-layers-to-conv-layers-8a43880a44ed

inputs = Input(shape = (256, 256, 3))

x = base_model(inputs, training = False)
x = Conv2D(filters = 512, kernel_size = (8, 8), activation = 'relu')(x)
x = Conv2D(filters = 128, kernel_size = (1, 1), activation = 'relu')(x)
x = tf.keras.layers.GlobalAveragePooling2D()(x)
output_ =Dense(16, activation='softmax')(x)

final_model_3 = Model(inputs = inputs, outputs = output_, name = 'Model_3')

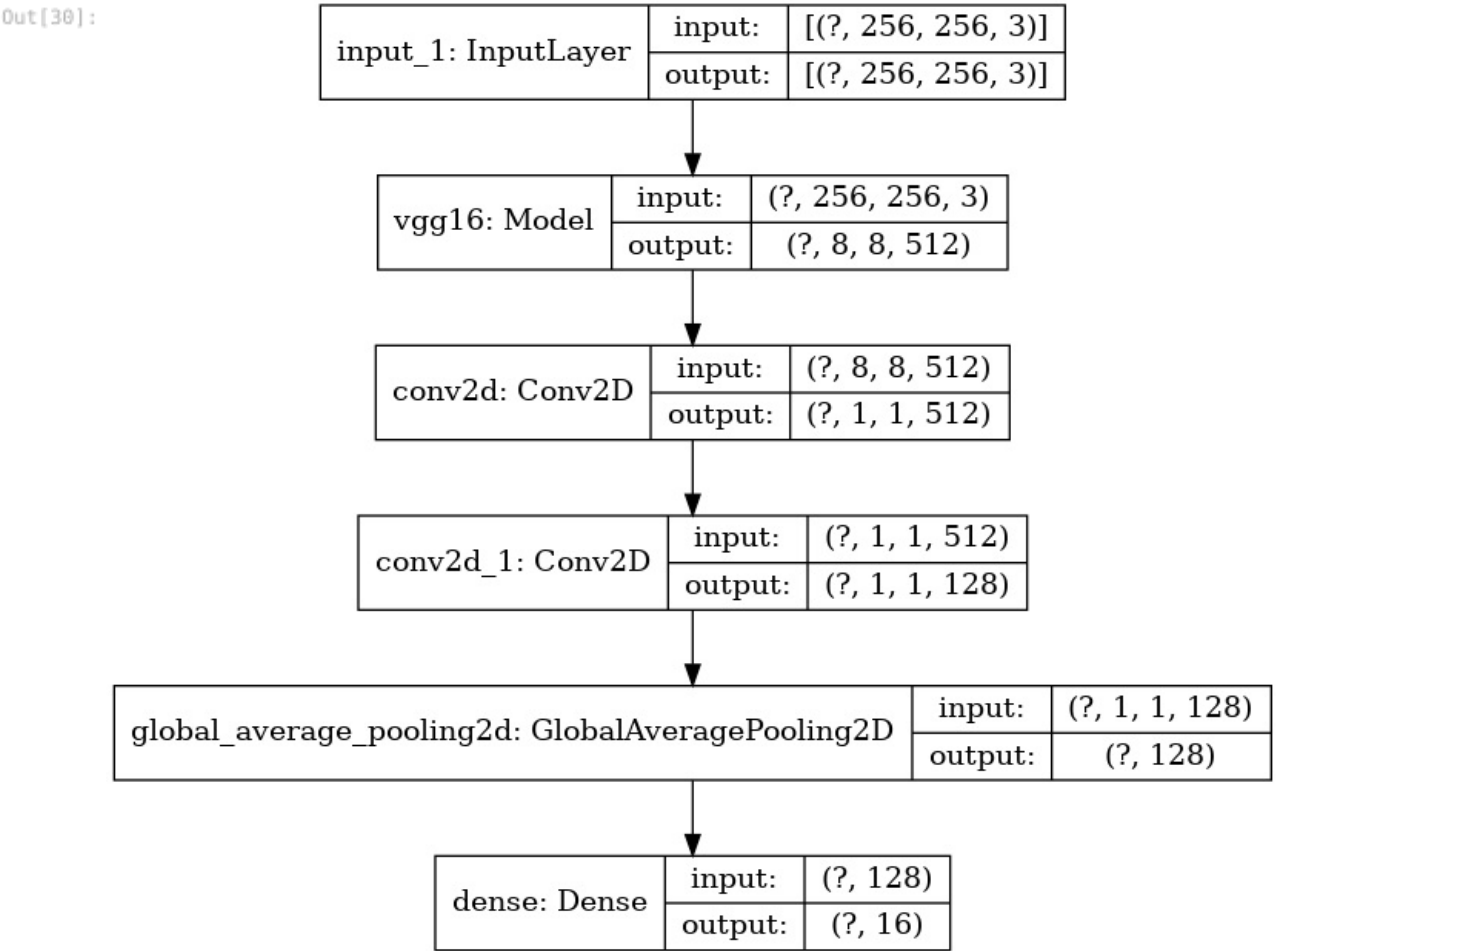
final_model_3.summary()
```

Model: "Model_3"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
vgg16 (Model)	(None, 8, 8, 512)	14714688
conv2d (Conv2D)	(None, 1, 1, 512)	16777728
conv2d_1 (Conv2D)	(None, 1, 1, 128)	65664
global_average_pooling2d (Gl	(None, 128)	0
dense (Dense)	(None, 16)	2064

=====
Total params: 31,560,144
Trainable params: 26,284,688
Non-trainable params: 5,275,456
=====

```
In [30]: plot_model(final_model_3, to_file = 'Model_Plots/Model_3.png', show_shapes = True)
```



```
In [31]: final_model_3.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])
```

```
In [32]: logdir = 'logs/model_3/' + datetime.now().strftime('%Y%m%d_%H%M%S')
tensorBoard = TensorBoard(log_dir = logdir, histogram_freq = 1)

filepath = 'modelCheck/model_3/epo_{epoch:02d}-accu_{val_accuracy:.4f}.hdf5'
modelCheck = ModelCheckpoint(filepath, monitor = 'val_accuracy', verbose = 1)

callBack_List = [tensorBoard, modelCheck, earlyStopping, reduceLr]
```

```
In [33]: EPOCHS_3 = 5
model_3 = final_model_3.fit(train_gen, steps_per_epoch = step_train, epochs = EPOCHS_3,
                             validation_data = test_gen, validation_steps = step_test, callbacks = callBack_List)
```


```
Epoch 1/5
1125/1125 [=====] - ETA: 0s - loss: 2.7965 - accuracy: 0.0618
Epoch 00001: saving model to modelCheck/model_3/epo_01-accu_0.0584.hdf5
1125/1125 [=====] - 683s 607ms/step - loss: 2.7965 - accuracy: 0.0618 - val_loss: 2.7728
- val_accuracy: 0.0584 - lr: 0.0010
Epoch 2/5
1125/1125 [=====] - ETA: 0s - loss: 2.7728 - accuracy: 0.0618
Epoch 00002: saving model to modelCheck/model_3/epo_02-accu_0.0587.hdf5
1125/1125 [=====] - 684s 608ms/step - loss: 2.7728 - accuracy: 0.0618 - val_loss: 2.7728
- val_accuracy: 0.0587 - lr: 0.0010
Epoch 3/5
1125/1125 [=====] - ETA: 0s - loss: 2.7728 - accuracy: 0.0629
Epoch 00003: saving model to modelCheck/model_3/epo_03-accu_0.0596.hdf5
1125/1125 [=====] - 686s 609ms/step - loss: 2.7728 - accuracy: 0.0629 - val_loss: 2.7730
- val_accuracy: 0.0596 - lr: 0.0010
Epoch 4/5
1125/1125 [=====] - ETA: 0s - loss: 2.7728 - accuracy: 0.0595
Epoch 00004: saving model to modelCheck/model_3/epo_04-accu_0.0630.hdf5
1125/1125 [=====] - 678s 603ms/step - loss: 2.7728 - accuracy: 0.0595 - val_loss: 2.7730
- val_accuracy: 0.0630 - lr: 0.0010
Epoch 5/5
1125/1125 [=====] - ETA: 0s - loss: 2.7728 - accuracy: 0.0613
Epoch 00005: saving model to modelCheck/model_3/epo_05-accu_0.0604.hdf5


Epoch 00005: ReduceLROnPlateau reducing learning rate to 0.0009000000427477062.
1125/1125 [=====] - 682s 606ms/step - loss: 2.7728 - accuracy: 0.0613 - val_loss: 2.7728
- val_accuracy: 0.0604 - lr: 0.0010
Epoch 00005: early stopping
```

```
In [34]: final_model_3.save_weights('Model_Weights/best_model_3.h5')
```

```
In [35]: # https://wellsr.com/python/seaborn-line-plot-data-visualization/
# https://seaborn.pydata.org/generated/seaborn.lineplot.html

epochs_ = list(range(len(model_3.history['accuracy'])))
plt.figure(figsize = (12, 8))
sns.lineplot(y = 'accuracy', data = model_3.history, x = epochs_, label = 'Accuracy',
              color = 'b', marker = 'X', mfc = 'red', ms = 8)
sns.lineplot(y = 'val_accuracy', data = model_3.history, x = epochs_,
              label = 'Validation Accuracy', color = 'g', marker = 'X', mfc = 'red', ms = 8)
plt.axhline(0.06, color = '#E44CF6', label = '6% Threshold')
plt.title('Model 3 : Accuracy vs Epochs') ; plt.legend(loc = 4)
plt.xticks(epochs_) ; plt.xlabel('Epochs') ; plt.ylabel('Accuracy') ; plt.show()
```





Observation

- The accuracy for this model is only less than 10%.
- Morethan 83% of parameters are trainable from total parameters.
- The loss was more than 2.5. This both indicates the model performance was pritty bad.
- This indicates that training VGG16 layaers won't give better accuracy always.
- Since from the grapgh indications, it is clear that may be with more epoch values may be we will get much more better results.
- Higher trainable parameters will give better result at higher epochs.

Models Summary

In [36]: `# https://pypi.org/project/prettytable/`

```
from prettytable import PrettyTable

table = PrettyTable()

table.field_names = ['Model', 'Number of epochs', 'Accuracy',
                    'Validation Accuracy']

table.add_row(['Model_1', EPOCHS_1, round(model_1.history['accuracy'][-1], 4),
              round(model_1.history['val_accuracy'][-1], 4)])

table.add_row(['Model_2', EPOCHS_2, round(model_2.history['accuracy'][-1], 4),
              round(model_2.history['val_accuracy'][-1], 4)])

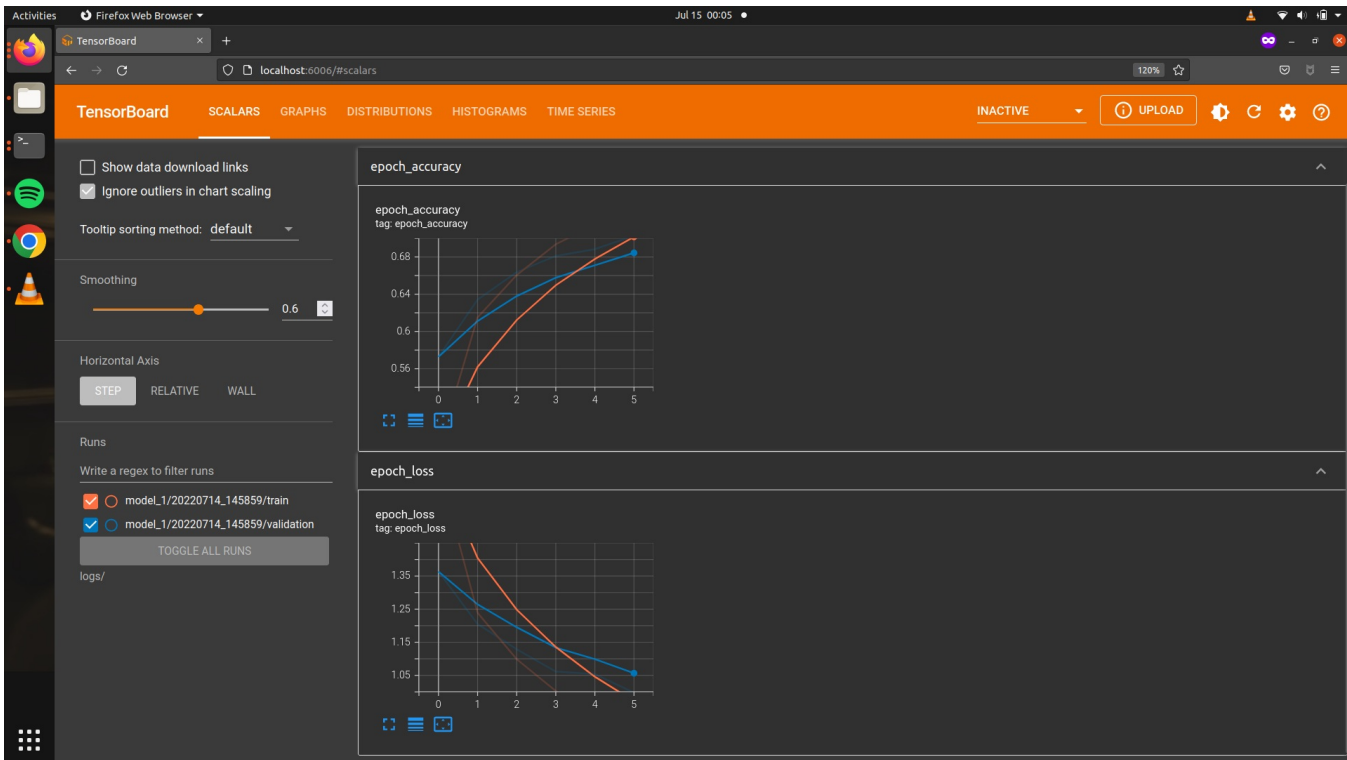
table.add_row(['Model_3', EPOCHS_3, round(model_3.history['accuracy'][-1], 4),
              round(model_3.history['val_accuracy'][-1], 4)])

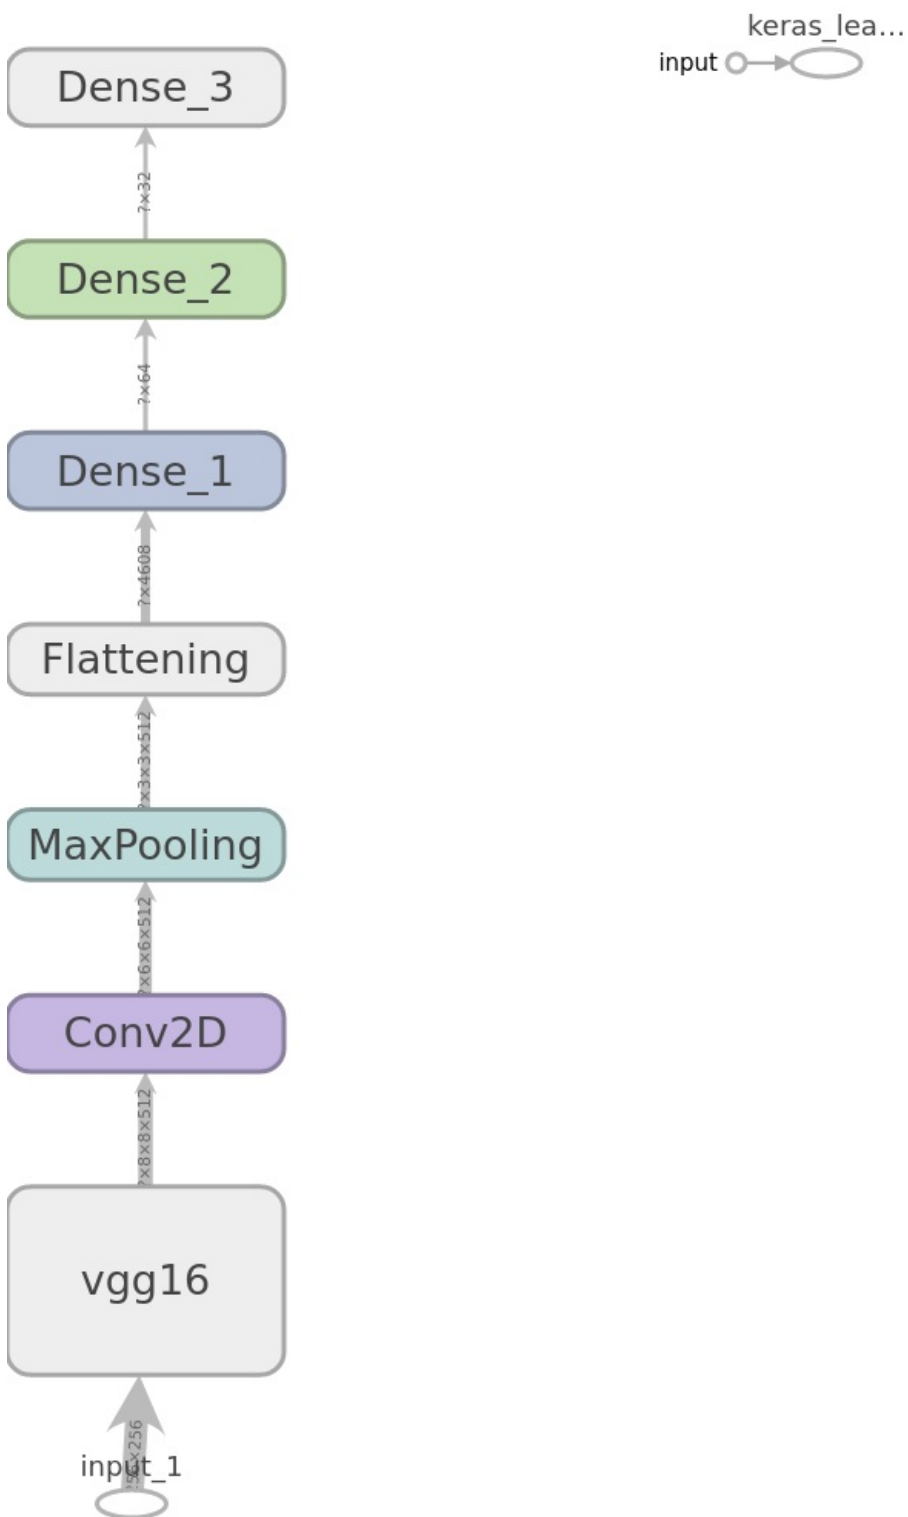
print(table)
```

Model	Number of epochs	Accuracy	Validation Accuracy
Model_1	6	0.7345	0.7028
Model_2	6	0.7603	0.7246
Model_3	5	0.0613	0.0604

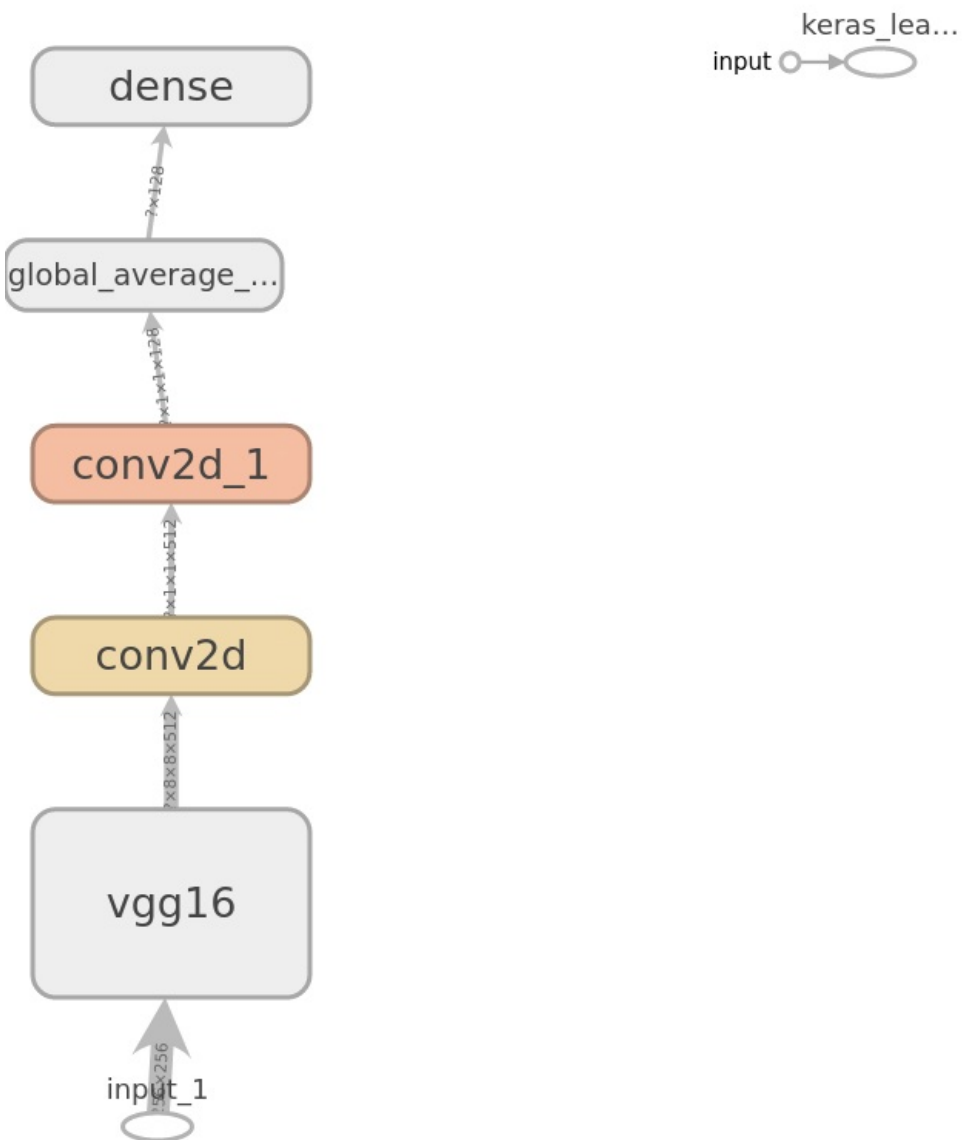
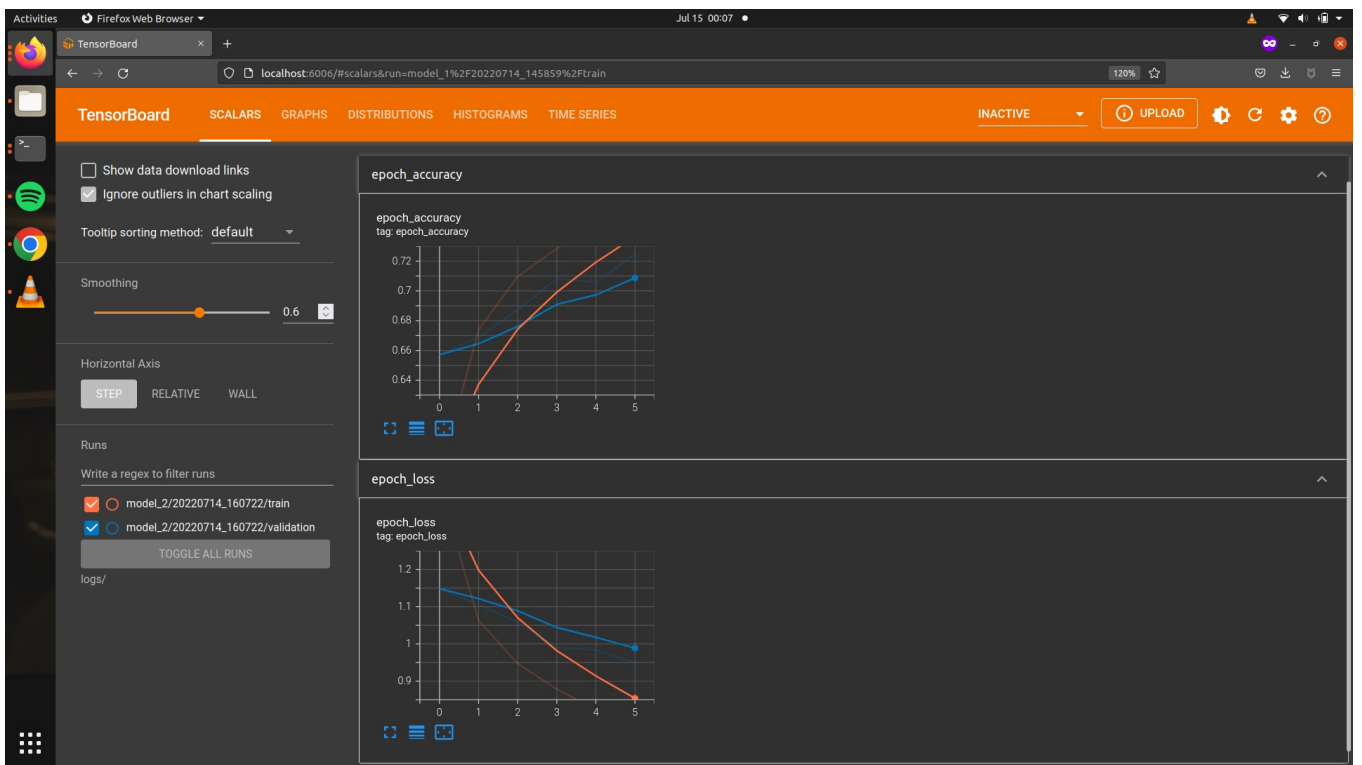
TensorBoard Outputs

Model - 1





Model - 2



Model - 3

